



Batch Size (B)	128
Optimizer	AdamW
Initial Learning Rate	1e-3
Weight Decay	1e-4
Scheduler	OneCycleLR

Loss Function	Focal ($\gamma=2.0$)
Label Smoothing	0.01
Grad Clipping	1.0

Hyperparameters & Config