


Training Loop Architecture



Hyperparameters

Epochs	30
Batch Size	128
Optimizer	AdamW
LR	1e-3
Weight Decay	1e-4
Scheduler	OneCycleLR
Loss	Focal ($\gamma=2$)
Grad Clip	1.0
Mixup α	0.1